

Few-shot Learning with Meta-Learning for Earth Observation

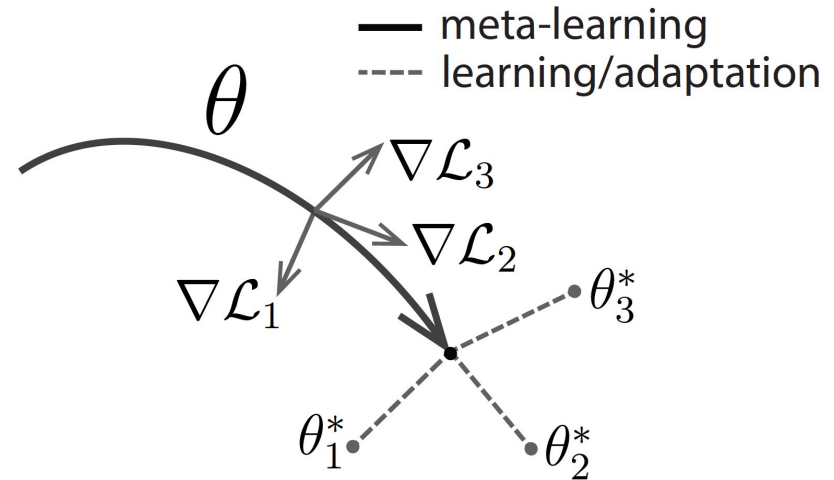
Cesar Luis Aybar Camacho

** Student at Copernicus Master in Digital Earth (EMCDE)*

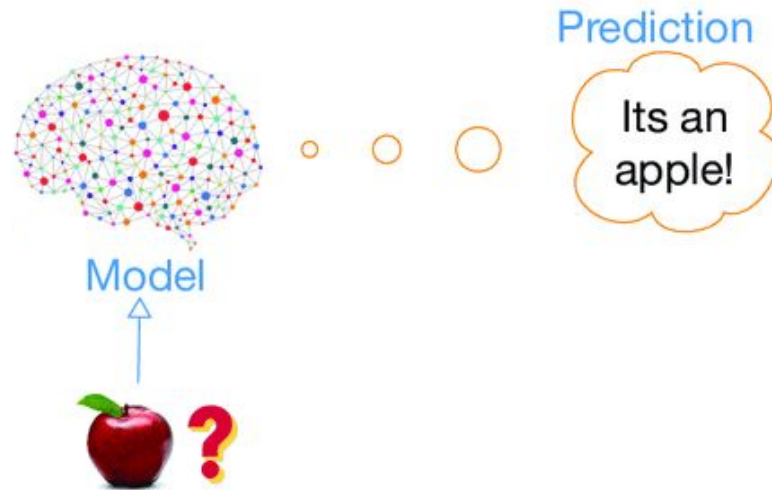
April 22, 2021



With the support of the
Erasmus+ Programme
of the European Union



Standard Supervised Learning Problem



Input data



Annotations

These are
apples



Model



Prediction

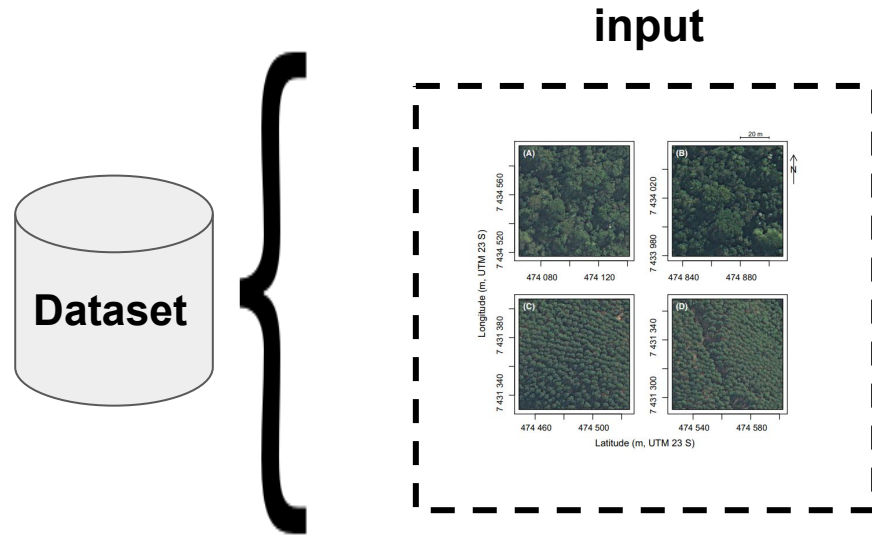
Its an
apple!

How we can map forest disturbance in the Atlantic rainforest?

[Fabien H. Wagne, et. al, 2019](#)



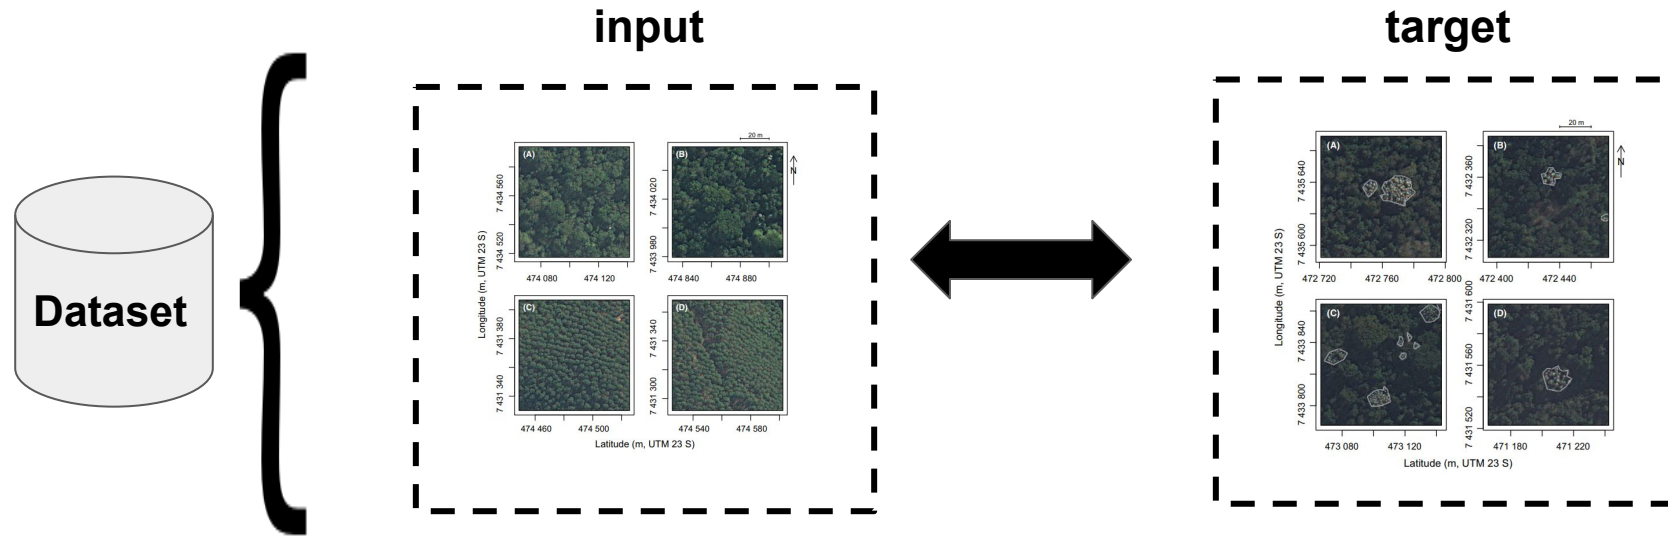
How we can map forest disturbance in the Atlantic rainforest?



[Fabien H. Wagne, et. al, 2019](#)



How we can map forest disturbance in the Atlantic rainforest?


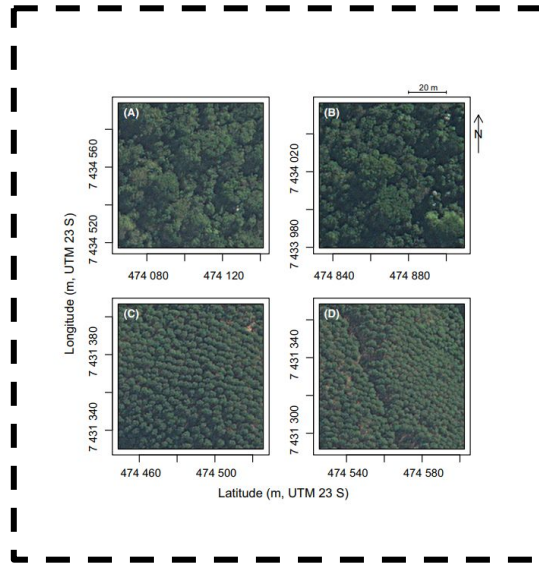


[Fabien H. Wagne, et. al, 2019](#)

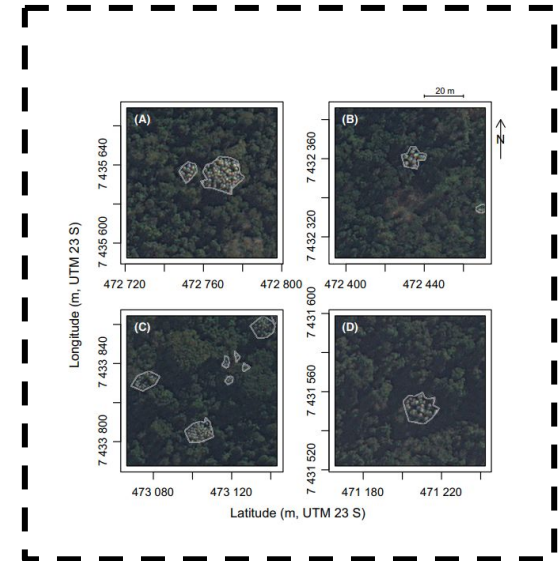


How we can map forest disturbance in the Atlantic rainforest?

input

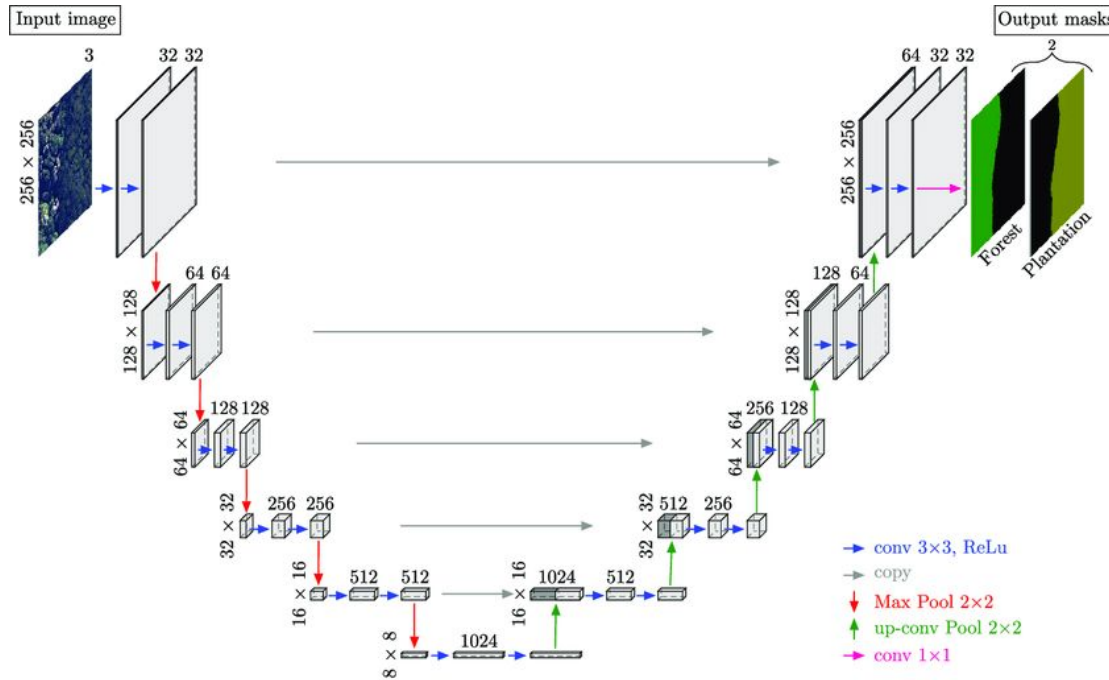


target



[Fabien H. Wagne, et. al. 2019](#)

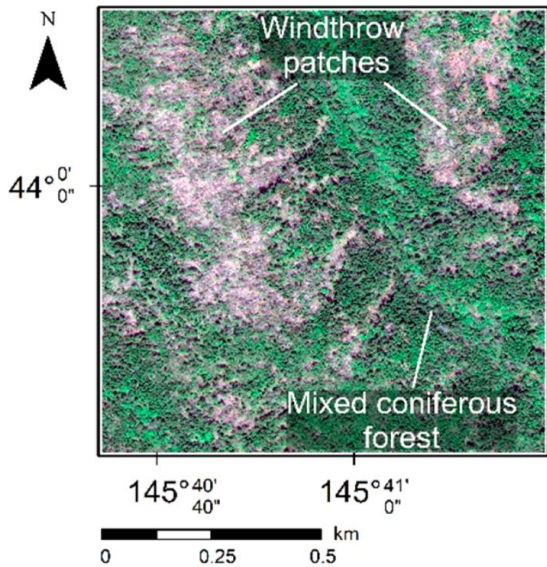





U-NET

[Fabien H. Wagne, et. al, 2019](#)





U-NET



[Fabien H. Wagne, et. al, 2019](#)



Shortcomings

- Large dataset are not always available.



Shortcomings


- **Large dataset are not always available.**
- **Need retrain every time new data is added.**



Shortcomings


- Large dataset are not always available.
- Need retrain every time new data is added.

U-NET



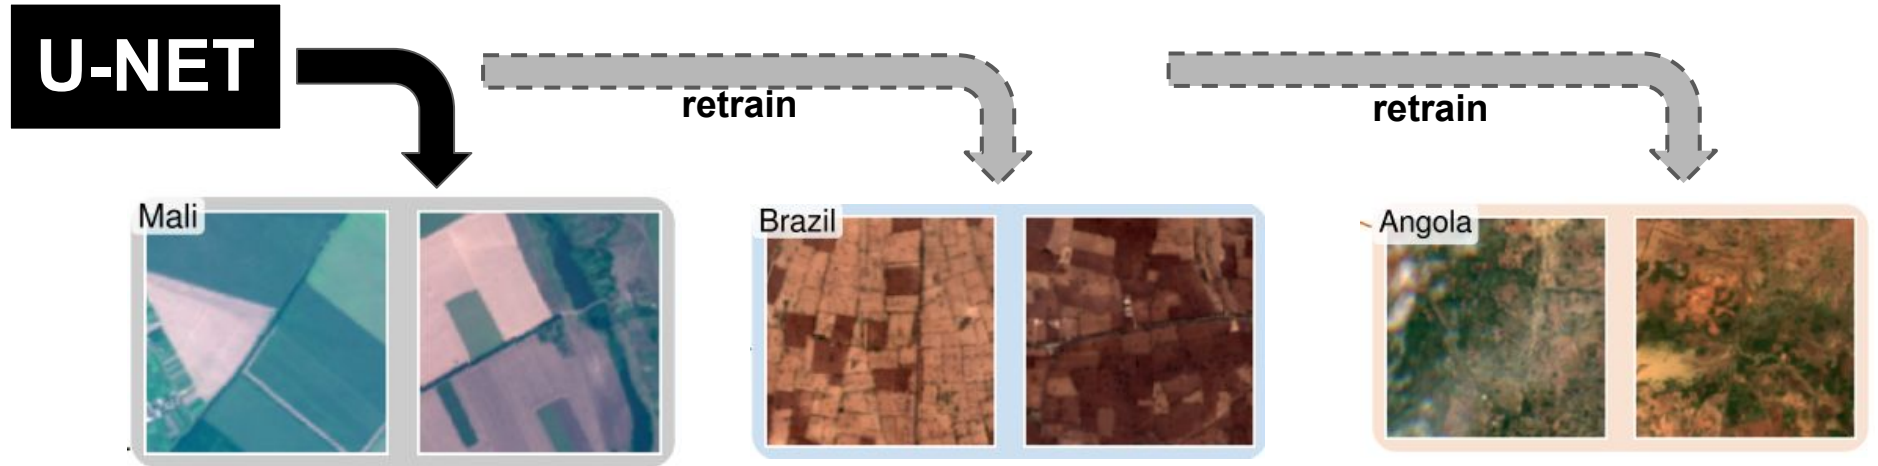
Shortcomings

- Large dataset are not always available.
- Need retrain every time new data is added.



Shortcomings


- Large dataset are not always available.
- Need retrain every time new data is added.



One model per region




One model on pooled data




Beyond Supervised Learning

Semantic segmentation approaches:


self-supervised




Supervised




Semi- Supervised



Weakly-supervised




few-shot learning




Semantic segmentation approaches:


self-supervised




Supervised




Semi- Supervised




Weakly-supervised



few-shot learning




Few-shot Learning 🔥



Tom




What is her name?



Tom






Tom





What is her name?



Tom



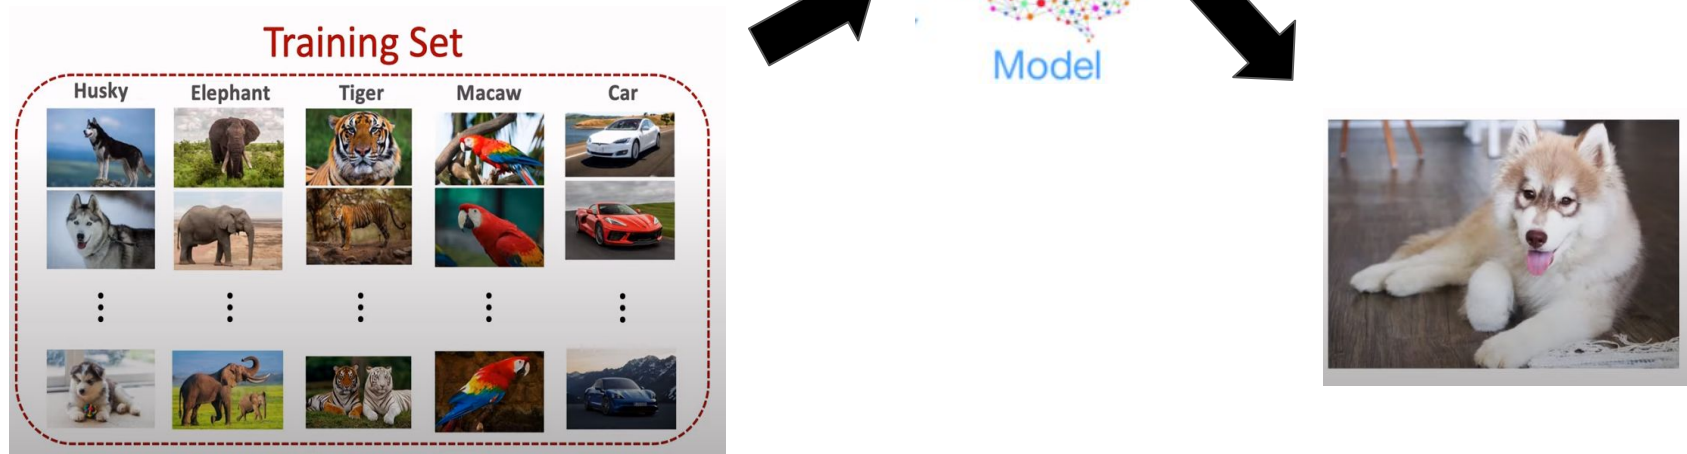
Support Dataset S_n

What is her name?



Query Dataset Q_n

Standard SL



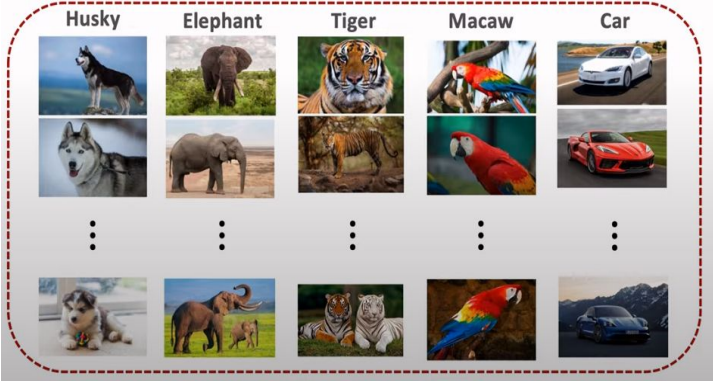


[Shusen Wang video](#)



Standard SL

Training Set




[Shusen Wang video](#)



Standard SL


Training Set



[Shusen Wang video](#)



Standard SL


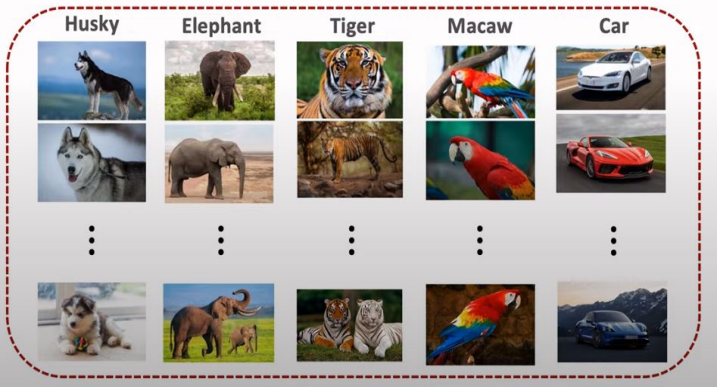


[Shusen Wang video](#)



Few Shot L

Training Set


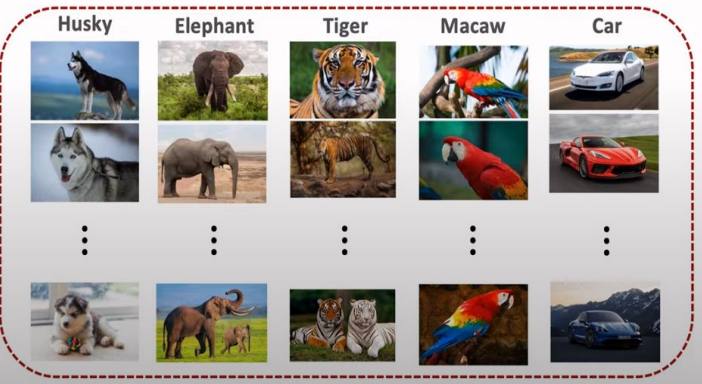


[Shusen Wang video](#)



Few Shot L

Training Set



Support Set:



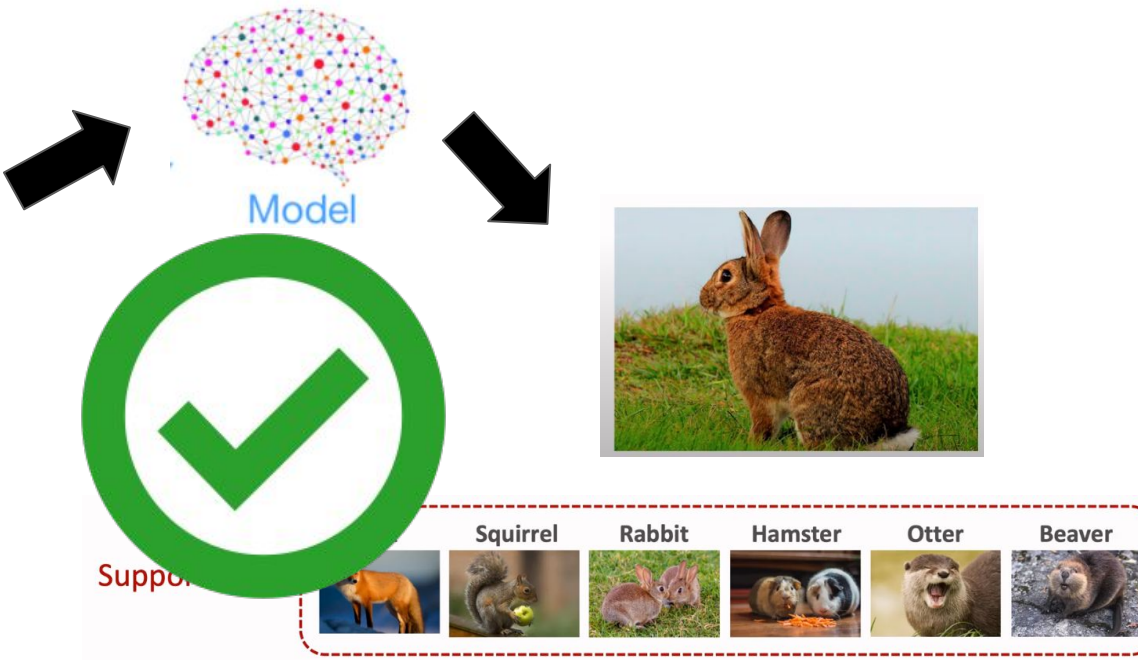
[Shusen Wang video](#)



Few Shot L

Training Set

Husky	Elephant	Tiger	Macaw	Car
				
				
⋮	⋮	⋮	⋮	⋮
				




[Shusen Wang video](#)



What is her name?

Basic Idea



1. Divide the **dataset** in **query** and **support**.
2. Learn a **similarity function**.
3. Apply the **similarity function** to the predictions.



Few shot Learning approaches:

- **Feature Transfer:** Standard transfer learning, Baseline++ (Chen et al. 2019), Simpleshot (Wang et al. 2019), etc.
- **Metric Learning:** Matching Networks (Vinyals et al. 2016), Prototypical Networks (Snell et al. 2017), Relation Networks (Sunget et al. 2018), etc.
- **Meta-learning:** **Model-Agnostic Meta-Learning (MAML, Finn et al. 2017), MAML ++ (Antoniou et al. 2019), Meta-SGD (Li et al. 2017), etc.**
- **Bayesian methods:** Bayesian MAML (Yoon et al. 2018), VERSA (Gordon et al. 2019), ALPaCA (Harrison et al. 2018), etc.

[Bayesian Meta-Learning for the Few-Shot Setting via Deep Kernels - Massimiliano Patacchiola 2020 - NeurIPS 2020](#)



MetaLearning + Few-Shot Learning


In meta-learning models learn how to learn!

Single Task (from scratch)




Learn task and perform task

Pretraining and fine-tuning



Refresh task of interest

meta-learning



Quickly learn a new task

[Bayesian Meta-Learning for the Few-Shot Setting via Deep Kernels - Massimiliano Patacchiola 2020 - NeurIPS 2020](#)



Model-Agnostic Meta-Learning

MAML

Agnostic, in the sense that the method can be used in different contexts, few-shot learning is a particular case.




Figure 1. Diagram of our model-agnostic meta-learning algorithm (MAML), which optimizes for a representation θ that can quickly adapt to new tasks.




MAML Intuition

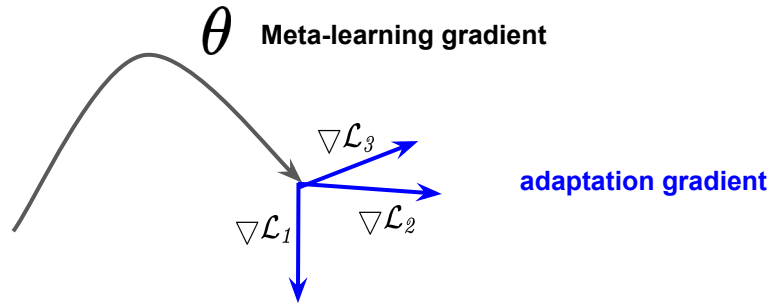
$$\min_{\theta} \theta - \alpha \nabla_{\theta} \mathcal{L}(\theta, \mathcal{D})$$



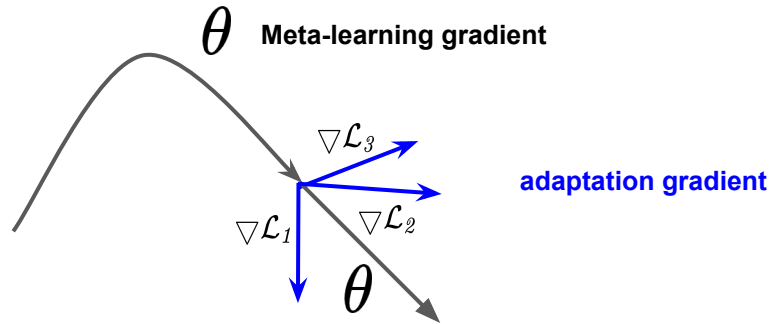
MAML Intuition



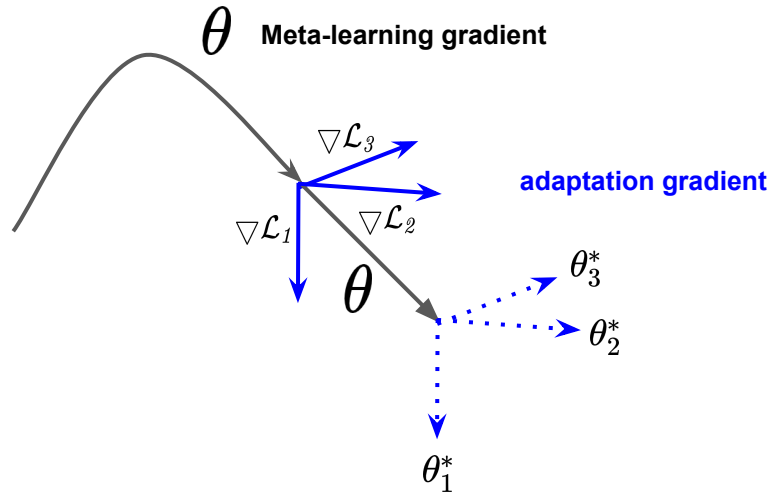
MAML Intuition



MAML Intuition



MAML Intuition




MAML Intuition

\mathcal{F}_θ



MAML Intuition


1. Copy Model per task



MAML Intuition

1. Copy Model per task

2. Support set train




MAML Intuition

1. Copy Model per task

2. Support set train

3. Calculate query set loss




MAML Intuition

1. Copy Model per task

2. Support set train

3. Calculate query set loss

4. Sum task losses




MAML Intuition

1. Copy Model per task

2. Support set train

3. Calculate query set loss

4. Sum task losses



MAML Gradient Descent

$$\min_{\theta} \theta - \alpha \nabla_{\theta} \mathcal{L}(\theta, \mathcal{D})$$



MAML Gradient Descent

The diagram shows the MAML Gradient Descent equation: $\min_{\theta} \theta - \alpha \nabla_{\theta} \mathcal{L}(\theta, \mathcal{D})$. Handwritten pink annotations identify the components: a wavy line under θ is labeled 'Updated parameter'; a bracket under $\mathcal{L}(\theta, \mathcal{D})$ is labeled 'Loss function'; a line from \mathcal{D} is labeled 'Data Points'; and a bracket under α is labeled 'step size hyperparameter'.

$$\min_{\theta} \theta - \alpha \nabla_{\theta} \mathcal{L}(\theta, \mathcal{D})$$

Updated parameter

Loss function

Data Points

step size hyperparameter



MAML Gradient Descent

$$\min_{\theta} \sum_{\tau_i \sim p(\tau)} \mathcal{L}(\theta - \alpha \nabla_{\theta} \mathcal{L}(\theta, \mathcal{D}_{\tau_i}^{\mathcal{S}}), \mathcal{D}_{\tau_i}^{\mathcal{Q}}) = \sum_{\tau_i \sim p(\tau)} \mathcal{L}(\theta', \mathcal{D}_{\tau_i}^{\mathcal{Q}})$$



MAML Gradient Descent

$$\min_{\theta} \sum_{\tau_i \sim p(\tau)} \mathcal{L}(\theta - \alpha \nabla_{\theta} \mathcal{L}(\theta, \mathcal{D}_{\tau_i}^{\mathcal{S}}), \mathcal{D}_{\tau_i}^{\mathcal{Q}}) = \sum_{\tau_i \sim p(\tau)} \mathcal{L}(\theta', \mathcal{D}_{\tau_i}^{\mathcal{Q}})$$

Diagram illustrating the MAML Gradient Descent equation with annotations:

- MetaModel parameter**: Points to θ .
- Task**: Points to $\tau_i \sim p(\tau)$.
- Loss function MetaLearning**: Points to the inner expression $\mathcal{L}(\theta - \alpha \nabla_{\theta} \mathcal{L}(\theta, \mathcal{D}_{\tau_i}^{\mathcal{S}}), \mathcal{D}_{\tau_i}^{\mathcal{Q}})$.
- Gradient Descent Adaptation**: Points to the update term $\theta - \alpha \nabla_{\theta} \mathcal{L}(\theta, \mathcal{D}_{\tau_i}^{\mathcal{S}})$.
- Support Set**: Points to $\mathcal{D}_{\tau_i}^{\mathcal{S}}$.
- Query Set**: Points to $\mathcal{D}_{\tau_i}^{\mathcal{Q}}$.
- Updated Parameters**: Points to θ' .



Model-Agnostic Meta-Learning

```
1 model = ConvolutionalNeuralNetwork(out_features=5) #we suppose a 5-way setting
2 meta_optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)
3
4 #[...] outer loop starts here, sample batch of tasks
5
6 for task in batch:
7     train_inputs, train_targets = task['support'] #input-output train pairs
8     test_inputs, test_targets = task['query'] #input-output test pairs
9
10    train_logit = model(train_input)
11    inner_loss = F.cross_entropy(train_logit, train_target) #on train set
12    model.zero_grad()
13    grads = torch.autograd.grad(inner_loss, model.meta_params(), create_graph=True)
14    params = OrderedDict()
15    for (name, param), grad in zip(model.meta_named_pars(), grads):
16        params[name] = param - step_size * grad
17    test_logit = model(test_input, params=params) #assign params to model
18    #Notice the `+=` which is used to accumulate the loss for each task
19    outer_loss += F.cross_entropy(test_logit, test_target) #on test set
20
21 outer_loss.backward()
22 meta_optimizer.step()
```

Algorithm 2 MAML for Few-Shot Supervised Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

- 1: randomly initialize θ
 - 2: **while** not done **do**
 - 3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$
 - 4: **for all** \mathcal{T}_i **do**
 - 5: Sample K datapoints $\mathcal{D} = \{\mathbf{x}^{(j)}, \mathbf{y}^{(j)}\}$ from \mathcal{T}_i
 - 6: Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ using \mathcal{D} and $\mathcal{L}_{\mathcal{T}_i}$ in Equation (2) or (3)
 - 7: Compute adapted parameters with gradient descent:
 $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$
 - 8: Sample datapoints $\mathcal{D}'_i = \{\mathbf{x}^{(j)}, \mathbf{y}^{(j)}\}$ from \mathcal{T}_i for the meta-update
 - 9: **end for**
 - 10: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$ using each \mathcal{D}'_i and $\mathcal{L}_{\mathcal{T}_i}$ in Equation 2 or 3
 - 11: **end while**
-



Model-Agnostic Meta-Learning

```
1 model = ConvolutionalNeuralNetwork(out_features=5) #we suppose a 5-way setting
2 meta_optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)
3
4 #[...] outer loop starts here, sample batch of tasks
5
6 for task in batch:
7     train_inputs, train_targets = task['support'] #input-output train pairs
8     test_inputs, test_targets = task['query'] #input-output test pairs
9
10    train_logit = model(train_input)
11    inner_loss = F.cross_entropy(train_logit, train_target) #on train set
12    model.zero_grad()
13    grads = torch.autograd.grad(inner_loss, model.meta_params(), create_graph=True)
14    params = OrderedDict()
15    for (name, param), grad in zip(model.meta_named_pars(), grads):
16        params[name] = param - step_size * grad
17    test_logit = model(test_input, params=params) #assign params to model
18    #Notice the `+=` which is used to accumulate the loss for each task
19    outer_loss += F.cross_entropy(test_logit, test_target) #on test set
20
21 outer_loss.backward()
22 meta_optimizer.step()
```

Algorithm 2 MAML for Few-Shot Supervised Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters


- 1: randomly initialize θ
 - 2: **while** not done **do**
 - 3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$
 - 4: **for all** \mathcal{T}_i **do**
 - 5: Sample K datapoints $\mathcal{D} = \{\mathbf{x}^{(j)}, \mathbf{y}^{(j)}\}$ from \mathcal{T}_i
 - 6: Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ using \mathcal{D} and $\mathcal{L}_{\mathcal{T}_i}$ in Equation (2) or (3)
 - 7: Compute adapted parameters with gradient descent:
 $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$
 - 8: Sample datapoints $\mathcal{D}'_i = \{\mathbf{x}^{(j)}, \mathbf{y}^{(j)}\}$ from \mathcal{T}_i for the meta-update
 - 9: **end for**
 - 10: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$ using each \mathcal{D}'_i and $\mathcal{L}_{\mathcal{T}_i}$ in Equation 2 or 3
 - 11: **end while**
-



MAML pro vs cons



- Elegant and neat.
- Fully differentiable method
- Agnostic (easily adapted to multiple setting).



- Unstable, hard to train.
- High order derivatives.
- Vanishing gradient.



HOW TO TRAIN YOUR MAML

Antreas Antoniou

University of Edinburgh

{a.antoniou}@sms.ed.ac.uk

Amos Storkey

University of Edinburgh


{a.storkey}@ed.ac.uk

Harrison Edwards

OpenAI, University of Edinburgh

{h.l.edwards}@sms.ed.ac.uk

Strided MAML vs Strided MAML++



- <https://paperswithcode.com/sota/few-shot-image-classification-on-mini-2>
- <https://paperswithcode.com/sota/few-shot-semantic-segmentation-on-fss-1000>



Earth Observation?

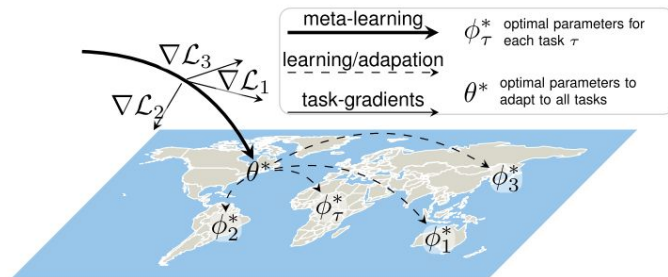
Meta-Learning for Few-Shot Land Cover Classification

Marc Rußwurm^{1,*,\dagger}, Sherrie Wang^{2,3,*}, Marco Körner¹, and David Lobell²

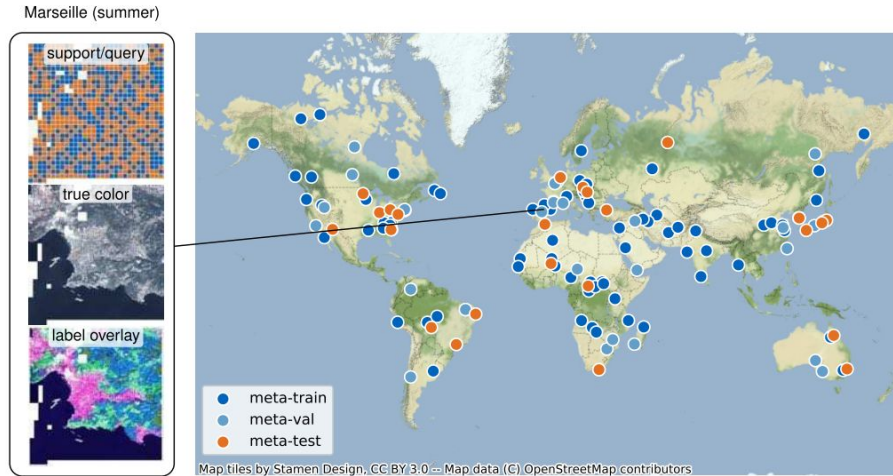
¹Technical University of Munich, Chair of Remote Sensing Technology

²Stanford University, Center on Food Security and the Environment

³Stanford University, Institute for Computational and Mathematical Engineering




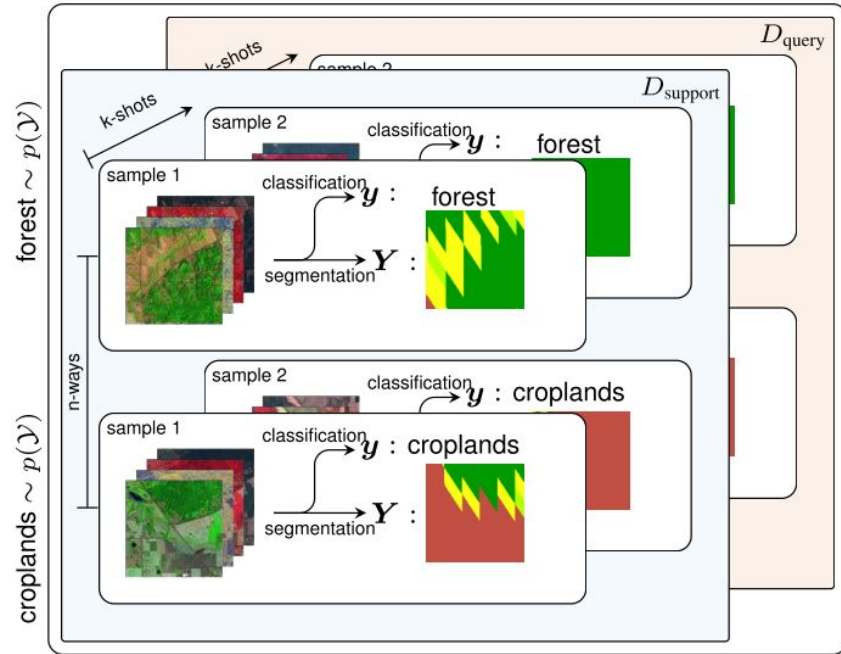
SEN12MS



- “Global” dataset
- Sentinel2 + MODIS
- 125 image tiles



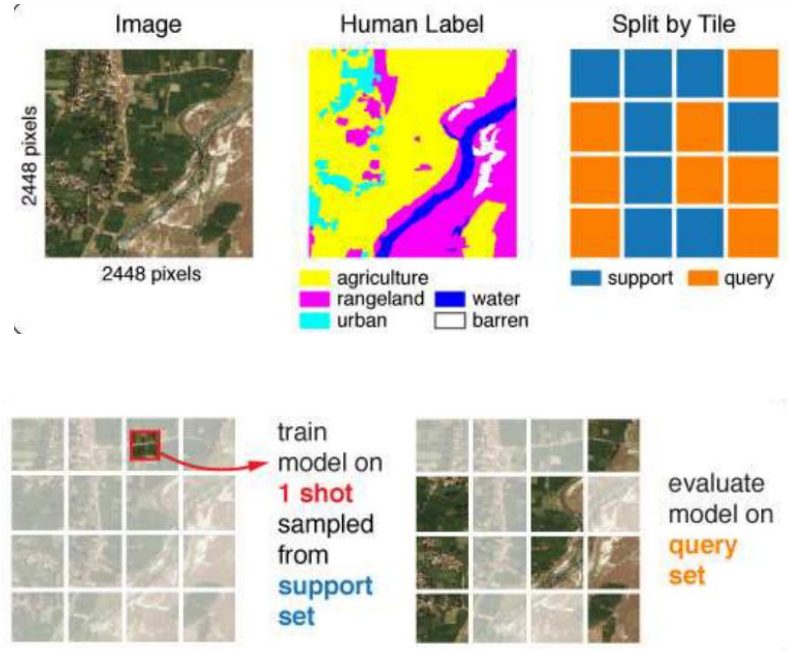
Geographic regions as meta-learning task



- **MAML adjusts** to new distribution in a single shot and outperforms baselines.



DeepGlobe




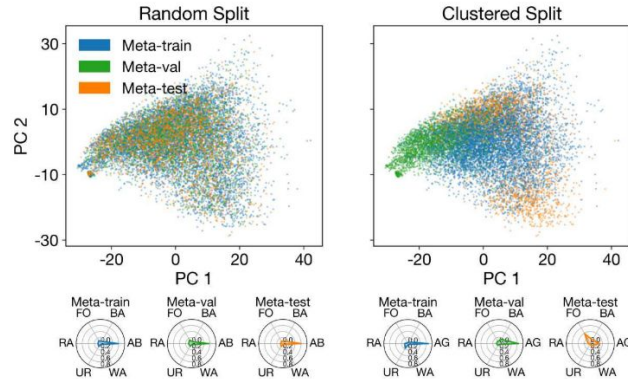

- RGB
- High Resolution (0.5 m)
- Semantic segmentation



When:

$$P_{train}(X, y) \neq P_{test}(X, y)$$


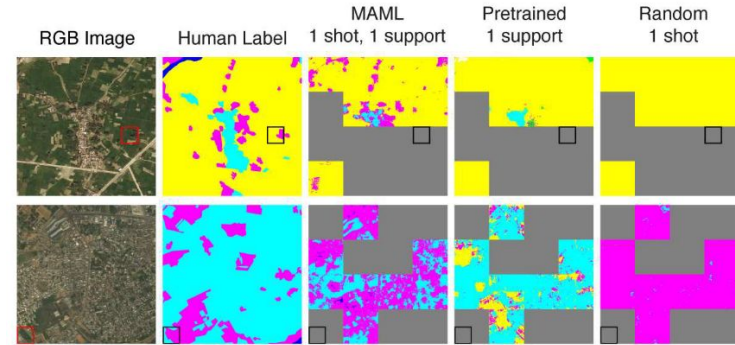
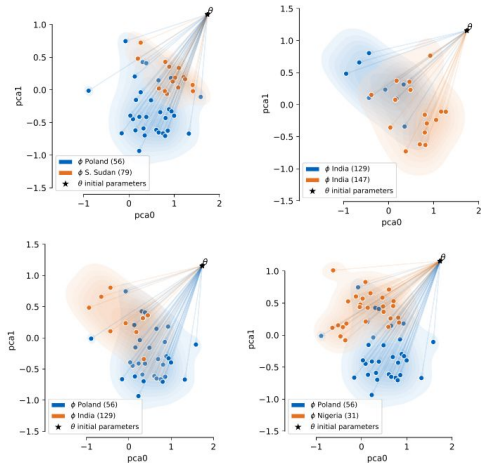
MAML outperforms pretraining



When:

$$P_{train}(X, y) \neq P_{test}(X, y)$$

MAML outperforms pretraining



Conclusion

Conclusions

- Results in computer vision paper show us that **meta-learning outperforms** pretraining and fine-tuning when the **meta-task tasks have data distribution that are different from meta-train tasks**.
- Current EO Deep Learning Dataset are a limitation.
- meta-learning framework can lead deep learning in Earth observation to a new era.



Muchas Gracias